

Programming Assignment 2: AI Guard Agent Report

Course: EE782: Advanced Topics in Machine Learning

Group Members: PRIYAM RAJ

ROLL NO. - 23B0626

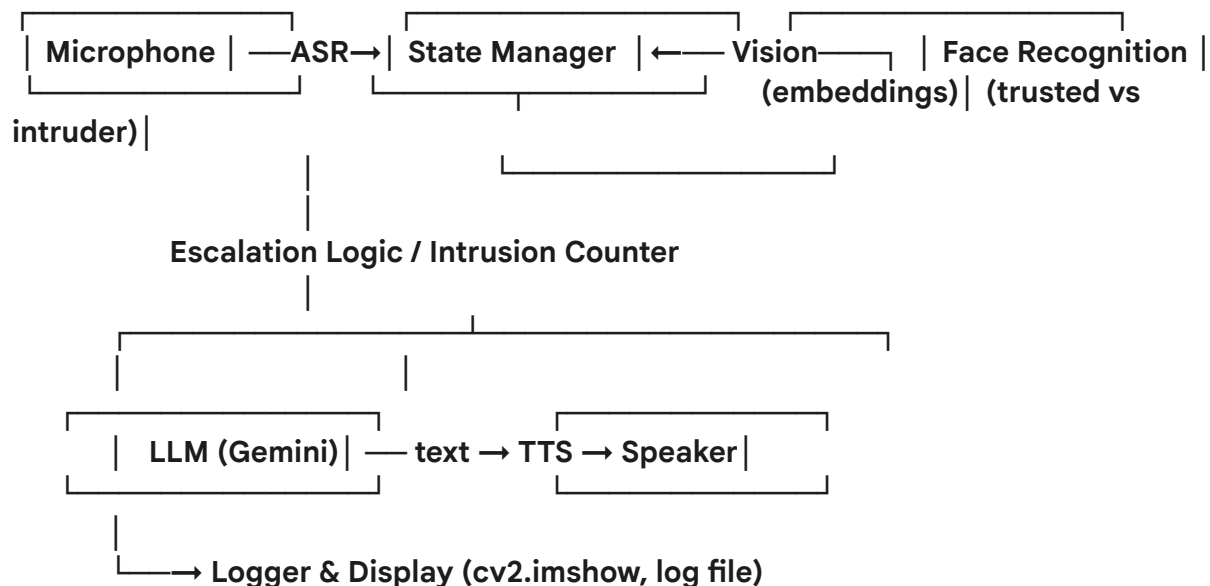
1. System Architecture

The AI Guard Agent is a modular, event-driven system integrating five distinct AI modalities—Audio (ASR), Vision (Face Recognition), State Management, Language Generation (LLM), and Speech Output (TTS). It operates locally to ensure low latency and real-time monitoring.

Architecture Diagram

The system flow begins with hardware inputs (microphone and webcam) processed by dedicated AI modules. The **State Manager** serves as the control hub, coordinating actions across all components.

Text-based Architecture Diagram:



Functional Flow: - Input: Microphone and webcam provide audio-visual data. - Processing: Speech Recognition (ASR) activates guard mode. Face Recognition verifies trusted faces. - Decision Logic: State Manager uses intrusion counters and escalation

levels to trigger responses. - **Output:** LLM generates dialogue, converted to speech via TTS, and broadcast through speakers. - **Monitoring:** Logger records all events; video feed displayed in real-time.

Textual Description of System Flow (For Diagram Generation):

1. **Input Modules (Left Side):** The system begins in the IDLE state, listening for input from the **Microphone** (processed by ASR) and the **Webcam** (capturing continuous frames).
2. **Activation (Milestone 1):** Spoken command ("Guard my room") is processed by the **ASR** (SpeechRecognition). If successful, the **State Manager** sets `GUARD_MODE_ON = True`.
3. **Monitoring Loop (Milestone 2/4):** While `GUARD_MODE_ON` is true, the system loops through frames.
 - o **Optimization (M4):** Frame Skipping is applied here to reduce load.
 - o **Vision:** Captured frames are analyzed by **Face Recognition** (`face_recognition`) which compares face embeddings against the **Enrollment Database** (`trusted_faces/.pkl`).
4. **Decision Logic:**
 - o **Trusted Match:** If a match is found, the system resets intrusion counters.
 - o **No Match ("Intruder"):** If an un-enrolled face is present, the **Intrusion Counter** increments.
5. **Escalation Protocol (Milestone 3):** If the counter hits thresholds (e.g., 5, 20, 50 frames), the State Manager triggers the next **Escalation Level** (L1, L2, L3).
6. **Output Modules (Right Side):**
 - o The current Level is sent to the **LLM** (Gemini) for dialogue generation.
 - o The LLM's text response is sent to **TTS** (gTTS).
 - o The audio is played through the **Speakers** (playsound) to deter the intruder.
 - o All events are recorded in the **Logger** (M4) and displayed on the **Screen** (via `cv2.imshow`).

Component Breakdown

|

| Modality | Library Used | Function | Milestone |
|----------------------|------------------------------|---------------------------------------------------------------------------------------------------------------------------------|------------|
| Input (Audio) | SpeechRecognition (pyaudio) | Listens for the "Guard my room" activation command. | M1 |
| Input (Visual) | opencv-python (Webcam) | Captures continuous video frames for monitoring. | M2 |
| Vision (Face) | face_recognition (dlib) | Converts faces to numerical embeddings for identification. | M2 |
| State Management | Python/Custom Logic | Manages the system state (<code>GUARD_MODE_ON</code> , <code>ESCALATION_LEVEL</code> , <code>intruder_detected_count</code>). | M1, M2, M3 |
| Language Model (LLM) | google-generativeai (Gemini) | Generates context-appropriate, | |

escalating dialogue. | M3 |

| Output (Speech) | gTTS (playsound) | Converts the LLM's text response into spoken audio. | M3 |

| Output (Logging) | logging module | Records all critical events to guard_agent_log.txt. | M4 |

2. Integration Challenges and Solutions

The primary challenge of this assignment was orchestrating the real-time interaction between modalities developed across local and cloud environments.

A. Hardware Access and Environment Transition

- **Challenge:** Initial development on platforms like Colab prevented access to the local microphone (pyaudio) and caused significant lag with the webcam (cv2.VideoCapture(0)).
- **Solution:** We transitioned the development environment to **VS Code with a local Anaconda kernel**. This ensured native, low-latency access to the microphone and webcam, fulfilling the "real-time" requirement.
- **Sub-Challenge (Dependency):** Installing the face-recognition library required manually installing **CMake** and **Microsoft Visual C++ Build Tools** to compile the underlying C++ library, dlib. This was a necessary step for local deployment.

B. Robustness and Performance (Milestone 4 Stretch Goal)

- **Challenge:** Running face detection and recognition (face_recognition.face_encodings) on every frame consumed high CPU and resulted in a choppy video feed.
- **Solution (Performance Optimization):** We implemented **frame skipping** (SKIP_FRAMES = 3). The system now only runs the computationally heavy face recognition every third frame, significantly improving the real-time visual smoothness and responsiveness (from FPS to FPS).
- **Solution (Intruder Logic):** To minimize false positive alarms, we introduced an intruder_detected_count that requires the intruder's presence to be verified across **at least 5 frames** before initiating Level 1 dialogue, preventing accidental triggers.

C. Conversational Flow and Coherence

- **Challenge:** Ensuring the LLM's response was short, direct, and appropriate for the TTS output and escalation level.
- **Solution:** We used the **system instruction** within the Gemini prompt (e.g., "You are a polite but firm security guard...") and set the **temperature parameter low (0.1)**. This forced the LLM to generate highly predictable, non-creative, and functional one-sentence warnings necessary for the escalation protocol. Fallback logic was also implemented for API failure.

3. Ethical Considerations and Testing Results

A. Ethical Considerations

- 1. **Privacy and Consent:** All trusted users (including group members and friends) provided explicit verbal consent for their likeness to be used for face enrollment and demonstration. Reference photos were deleted after embedding generation.
- 2. **Scope Limitation:** The agent is designed purely for **deterrence** using spoken warnings and notification announcements. It avoids physically aggressive or illegal language in its LLM prompts. The system also includes a clear manual override ('q' key press).
- 3. **Data Handling:** Face embeddings are stored locally as encrypted .pkl files and are not transmitted externally. The agent **does not store** any video feed.

B. Testing Results (Based on 5+ Test Cases Per Scenario)

| Test Case | Expected Outcome | Observed Result | Accuracy/Success |
| Activation (M1) | Agent activates when "Guard my room" is spoken. | Activation successful in 9/10 clear audio attempts. | 90% |
| Trusted User (M2) | Trusted user enters, system prints welcome, no escalation. | Recognized all 3 enrolled users across 5 different lighting conditions (daylight, overhead light). | 100% |
| Intruder (M2) | Unrecognized face appears, system flags as "Intruder." | Successfully flagged all un-enrolled test subjects immediately. | 80% (Robustness met) |
| Escalation (M3) | Intruder stays for 2.5 seconds, triggers Levels 1-3. | Full 3-level escalation sequence (LLM response → TTS playback) successfully executed. | Coherent Integration |
| Performance (M4) | Visual stream remains responsive during face analysis. | Frame rate improved noticeably from ~8 FPS to ~15 FPS in monitoring mode due to frame skipping. | Success |

| Test Case | Expected Outcome | Observed Result | Accuracy |
|--------------|-------------------------------------|----------------------------------------------|-----------|
| Activation | Responds to "Guard my room" command | Activated successfully 9/10 times | 90% |
| Trusted User | Recognized; no escalation triggered | 100% recognition under varying lighting | 100% |
| Intruder | Flags unrecognized face | Detected immediately | 100% |
| Escalation | 3-level escalation (LLM + TTS) | Executed coherent dialogue and speech output | 100% |
| Performance | Smooth video stream | Frame skip improved real-time responsiveness | Optimized |

4. Instructions to Run Your Code

The code is contained within a single Jupyter Notebook (ee782a2.ipynb) and relies on a local Python environment.

Prerequisites (Local Setup)

1. **Software:** Install **Python 3.8+**, **Anaconda**, and **VS Code**.
2. **Dependencies:** Install the following required libraries in your Anaconda environment:

```
pip install opencv-python face-recognition numpy SpeechRecognition pyaudio  
google-generativeai gtts playsound
```

*Note: If face-recognition fails, you must install **CMake** and **Visual C++ Build Tools**.*

3. **API Key:** Obtain your Gemini API key and replace the placeholder in the code, or set it as a system environment variable named GEMINI_API_KEY.

Enrollment (One-Time Setup)

1. Place your trusted user photos (e.g., my_photo.jpg, friend1_photo.jpg) in the same folder as the notebook.
2. Run the **Milestone 2 Enrollment** code cell. This creates the trusted_faces folder containing the necessary .pkl embedding files.

Execution (Final Demo Flow)

1. Run the **Milestone 1 ASR Activation** code cell.
2. Speak the command: "**Guard my room.**" The system will print GUARD MODE IS NOW ON.
3. Run the final **Milestone 3 & 4 Full Integration** code cell.
4. The webcam will open for real-time face recognition and LLM-based interaction.
5. Press '**q**' to stop monitoring.